

## Part 1:

The `ProjectCode1.py` has been modified to include a `--buffer_size` flag to specify the size of the finite buffer and the `--time_distribution` flag to specify the distribution for inter-arrival time and service time. For both questions in Part 1, I used a simulation time of 1000000 iterations.

### Q1:

I experimented with arrival rates of [0.1, 0.2, 0.4, 0.8], buffer sizes of [5, 10, 20, 40] given in packets, a  $\mu$  of 2 packets per second, and Poisson distribution for inter-arrival time and service time with the following commands:

```
python ProjectCode1.py --arrival_rates=0.1,0.2,0.4,0.8 --  
buffer_size=5 --time_distribution=poisson  
python ProjectCode1.py --arrival_rates=0.1,0.2,0.4,0.8 --  
buffer_size=10 --time_distribution=poisson  
python ProjectCode1.py --arrival_rates=0.1,0.2,0.4,0.8 --  
buffer_size=20 --time_distribution=poisson  
python ProjectCode1.py --arrival_rates=0.1,0.2,0.4,0.8 --  
buffer_size=40 --time_distribution=poisson
```

The results are presented in the tables below.

Arrival Rate ( $\lambda$ )	Buffer Size (packets)	Total Number of Packets	Number of Dropped Packets	Utilization
0.100	5	100516	33	0.201
0.200	5	200424	1255	0.398
0.400	5	400591	35942	0.729
0.800	5	799281	318417	0.962
0.100	10	100144	1	0.200
0.200	10	200479	25	0.401
0.400	10	401140	9722	0.784
0.800	10	799523	301425	0.997
0.100	20	100144	1	0.200
0.200	20	199805	3	0.398
0.400	20	400967	993	0.801
0.800	20	799964	300425	1.000
0.100	40	100144	1	0.200
0.200	40	199805	3	0.398
0.400	40	400672	5	0.801
0.800	40	799486	299468	1.000

## Discussion:

As I increased the arrival rate ( $\lambda$ ), I observed increases in total number of packets, number of dropped packets, and utilization.

When buffer size was 5, I observed a sublinear increase in utilization, e.g. when I doubled the arrival rate the utilization increased less than a factor of 2. However, as the buffer size increased from 5 to 10, 20, or 40, I observed a roughly linear increase in utilization.

I observed an exponential increase in number of dropped packets as I increased the arrival rate across all buffer sizes, whereas the total number of packets increased in a roughly-uniform manner.

The loss probability of packets increased superlinearly as the arrival rate increased across all buffer sizes, starting from a minimum of 0.001% at  $\lambda=0.100$  to a maximum of 37.701% at  $\lambda=0.800$ . Moreover, I observed that number of dropped packets and the loss probability was substantially higher for mid-level arrival rates (e.g.  $\lambda \in \{0.200, 0.400\}$ ) when buffer size was lower. Overall, we can conclude that buffer size is inversely proportional to the loss probability and arrival rate is proportional to the loss probability.

## Q2:

I experimented with arrival rates of [0.1, 0.2, 0.4, 0.8, 0.9, 0.99], buffer sizes of [5, 10, 20, 40], a  $\mu$  ( $\mu$ ) of 2 packets per second, and deterministic (constant) inter-arrival time and service time with the following commands:

```
python ProjectCode1.py --arrival_rates=0.1,0.2,0.4,0.8 --  
buffer_size=5 --time_distribution=constant  
python ProjectCode1.py --arrival_rates=0.1,0.2,0.4,0.8 --  
buffer_size=10 --time_distribution=constant  
python ProjectCode1.py --arrival_rates=0.1,0.2,0.4,0.8 --  
buffer_size=20 --time_distribution=constant  
python ProjectCode1.py --arrival_rates=0.1,0.2,0.4,0.8 --  
buffer_size=40 --time_distribution=constant
```

The results are presented in the table below.

Arrival Rate ( $\lambda$ )	Buffer Size (packets)	Total Number of Packets	Number of Dropped Packets	Utilization
0.100	5	99999	0	0.200
0.200	5	199999	0	0.400
0.400	5	399999	0	0.800
0.800	5	799999	30000	1.000
0.100	10	99999	0	0.200
0.200	10	199999	0	0.400

0.400	10	399999	0	0.800
0.800	10	799999	30000	1.000
0.100	20	99999	0	0.200
0.200	20	199999	0	0.400
0.400	20	399999	0	0.800
0.800	20	799999	30000	1.000
0.100	40	99999	0	0.200
0.200	40	199999	0	0.400
0.400	40	399999	0	0.800
0.800	40	799999	30000	1.000

### Discussion:

When the inter-arrival time and service time is made deterministic (constant) instead of Poisson distribution, we observe that the utilization has a perfectly linear, proportional relationship with the arrival rate independent of the buffer size. Furthermore, for the values we have experimented with, we observe that buffer size no longer has any impact on the number of dropped packets and the loss probability of packets: whereas arrival rates of 0.100, 0.200, and 0.400 yield 0 dropped packets and thus 0% loss, an arrival rate of 0.800 yields 30000 dropped packets and thus 3.75% loss probability. This is much lower compared to the probabilities reported in Q1, but nevertheless we still observe an proportional relationship between arrival rate and number of dropped packets and loss probability.

### Part 2

The `ProjectCode2.py` has been modified to make the simulation look like the topology given in Figure 1 and to match the initial assumptions and parameters given in the project specification.

### Q1:

I experimented with buffer sizes of [100, 250, 500, 1000, 10000, 120000] given in bytes, port rates of [100, 500, 1000, 2000] given in packets per second, and simulation time of 4000 iterations with the following commands:

```
python ProjectCode2.py --port_rate=100 --buffer_size=100
python ProjectCode2.py --port_rate=100 --buffer_size=250
python ProjectCode2.py --port_rate=100 --buffer_size=500
python ProjectCode2.py --port_rate=100 --buffer_size=1000
python ProjectCode2.py --port_rate=100 --buffer_size=10000
python ProjectCode2.py --port_rate=100 --buffer_size=120000
python ProjectCode2.py --port_rate=500 --buffer_size=100
python ProjectCode2.py --port_rate=500 --buffer_size=250
python ProjectCode2.py --port_rate=500 --buffer_size=500
python ProjectCode2.py --port_rate=500 --buffer_size=1000
python ProjectCode2.py --port_rate=500 --buffer_size=10000
```

```
python ProjectCode2.py --port_rate=500 --buffer_size=120000
python ProjectCode2.py --port_rate=1000 --buffer_size=100
python ProjectCode2.py --port_rate=1000 --buffer_size=250
python ProjectCode2.py --port_rate=1000 --buffer_size=500
python ProjectCode2.py --port_rate=1000 --buffer_size=1000
python ProjectCode2.py --port_rate=1000 --buffer_size=10000
python ProjectCode2.py --port_rate=1000 --buffer_size=120000
python ProjectCode2.py --port_rate=2000 --buffer_size=100
python ProjectCode2.py --port_rate=2000 --buffer_size=250
python ProjectCode2.py --port_rate=2000 --buffer_size=500
python ProjectCode2.py --port_rate=2000 --buffer_size=1000
python ProjectCode2.py --port_rate=2000 --buffer_size=10000
python ProjectCode2.py --port_rate=2000 --buffer_size=120000
```

The results are presented in the table below.

Port Rate	Buffer Size (bytes)	Average Wait Source 1 to Output 3	Average Wait Source 2 to Output 4	Packets Sent	Packets Received	Packets Dropped
100	100	0.01234	0.00422	6090	3383	2707
100	250	0.02332	0.00759	6157	4976	1181
100	500	0.02873	0.00958	6012	5357	655
100	1000	0.03110	0.01025	6122	5441	681
100	10000	0.03110	0.01025	6122	5441	681
100	120000	0.03110	0.01025	6122	5441	681
500	100	0.00251	0.00085	5944	3479	2465
500	250	0.00460	0.00151	6048	4872	1176
500	500	0.00585	0.00194	6084	5338	746
500	1000	0.00608	0.00198	6034	5423	611
500	10000	0.00608	0.00198	6034	5423	611
500	120000	0.00608	0.00198	6034	5423	611
1000	100	0.00121	0.00043	6039	3430	2609
1000	250	0.00231	0.00078	6024	4956	1068
1000	500	0.00291	0.00094	5983	5348	635
1000	1000	0.00298	0.00096	6152	5387	765
1000	10000	0.00298	0.00096	6152	5387	765
1000	120000	0.00298	0.00096	6152	5387	765
2000	100	0.00061	0.00021	6055	3395	2660
2000	250	0.00113	0.00038	6112	4933	1179
2000	500	0.00147	0.00049	6038	5405	633
2000	1000	0.00150	0.00050	6116	5362	754
2000	10000	0.00150	0.00050	6116	5363	753
2000	120000	0.00150	0.00050	6116	5363	753

### **Discussion:**

Buffer size, given in bytes as opposed to in packets as in the previous part, was proportional to both recorded average waits (i.e. one from source 1 to output 3 and one from source 2 to output 4); increasing the buffer size across all port rates increased the average wait duration. Across all port rates, we saw a saturation in average wait times after a buffer size of 1000. Port rate was inversely proportional to both average waits: a higher port rate corresponded with a lower average wait.

Buffer size was proportional to the number of packets received and inversely proportional to the number of packets dropped: a higher buffer size corresponded more packets received, less packets dropped, and hence a lower loss probability of packets. On the other hand, I observed that port rate did not have a visible impact on packets received, packets dropped, and hence the loss probability of packets.